CME 241: Reinforcement Learning for Stochastic Control Problems in Finance

Ashwin Rao

ICME, Stanford University

Meet your Instructor

- My educational background: Algorithms Theory & Abstract Algebra
- 10 years at Goldman Sachs (NY) *Rates/Mortgage Derivatives* Trading
- 4 years at Morgan Stanley as *Managing Director - Market Modeling*
- Founded Tech Startup *ZLemma*, Acquired by hired.com in 2015
- One of our products was algorithmic jobs/career guidance for students
- I've been teaching short/medium-length courses for 25 years
- Topics across Pure & Applied Math, CS, Programming, Finance
- Current Interest: A.I. for Dynamic Decisioning under Uncertainty
- V.P. Data Science at Target focused on Optimal Inventory Control
- Joined Stanford ICME as Adjunct in Fall 2018
- Apart from CME 241, I am a technical mentor to ICME students

Requirements and Setup

- (Light) Pre-requisites:
  - Undergraduate-level background in Applied Mathematics (Linear Algebra, Probability Theory, Optimization)
  - Background in Data Structures & Algorithms, with programming experience in numpy/scipy
  - Basic background in Pricing and Portfolio Theory, but we will do an overview of the requisite Finance/Economics
  - No background required in MDP, DP, RL (we will cover these topics from scratch)
- Install Python 3 and supporting editor/IDE (eg: PyCharm)
- Create git repo for this course (for assignments/sharing)
- Send the git repo details to the Course Assistant (for reviews/grading)
- Install LaTeX and supporting editor (eg: TeXShop)

Housekeeping

- Grade based on:
  - 20% Mid-Term Exam
  - 35% Final Exam
  - 35% Programming and Technical Writing Assignments
  - 10% Class Participation
- Passing grade fetches 3 credits, can be applied towards MCF degree
- Wed and Fri 4:30pm - 5:50pm, 01/07/2019 - 03/15/2019
- Classes in Building 200 (Lane History Corner) - Room 203
- Appointments: Any time Fridays or an hour before class Wednesdays
- Use appointments time to discuss theory as well as your code
- Course Web Site: cme241.stanford.edu
- Course Assistant (CA) is Jeffrey Gu (ICME student)
- My e-mail: ashwin.rao@stanford.edu, CA e-mail: jeffgu@stanford.edu

Resources

- I recommend Sutton-Barto as the companion book for this course
  - I won't follow the structure of Sutton-Barto book
  - But I will follow his approach/treatment
- I will follow the structure of David Silver's RL course
  - I encourage you to augment my lectures with David's lecture videos
  - Occasionally, I will wear away or speed up/slow down from this flow
- We will do a bit more Theory & a lot more coding (relative to above)
- You can freely use my open-source code for your coding work
  - I expect you to duplicate the functionality of above code in this course
- We will go over some classical papers on the Finance applications
- To understand in-depth the analytical solutions in simple settings
- I will augment the above content with many of my own slides
- All of this will be organized on the course web site

Learning Tenets (to thrive in this course)

- When learning Theory, blend notational rigor with concepts intuition
- LaTeX the mathematical formalism of lecture contents with precision
- Programming is a powerful tool to grasp mathematical concepts
- Code modularity/re-use is important (leverage OOP/FP paradigms)
- Always start with a simple model/algorithm, evolve incrementally
- Important to model various frictions encountered in real-world trading
- Discuss the ideas you learn with others (incl. with the CA and me)
- Remember this is an intro course. Avoid advanced RL at this stage
- Success := **Blending Theory, Algos/Code, Real-World Modeling**

A.I. for Dynamic Decisioning under Uncertainty

- Let's browse some terms used to characterize this branch of A.I.
- *Stochastic*: Uncertainty in key quantities, evolving over time
- *Optimization*: A well-defined metric to be maximized ("The Goal")
- *Dynamic*: Decisions need to a function of the changing situations
- *Control*: Overpower uncertainty by persistent steering towards goal
- Jargon overload due to confluence of Control Theory, O.R. and A.I.
- For language clarity, let's just refer to this area as *Stochastic Control*
- The core framework is called *Markov Decision Processes* (MDP)

Components of the MDP Framework

- The *Agent* and the *Environment* interact in a time-sequenced loop
- *Agent* responds to [*State*, *Reward*] by taking an *Action*
- *Environment* responds by producing next step's (random) *State*
- *Environment* also produces a (random) scalar denoted as *Reward*
- Each *State* is assumed to have the *Markov Property*, meaning:
  - Next *State/Reward* depends only on Current *State* (for a given *Action*)
  - Current *State* captures all relevant information from *History*
  - Current *State* is a sufficient statistic of the future (for a given *Action*)
- Goal of *Agent* is to maximize *Expected Sum* of all future *Reward*s
- By controlling the (*Policy* : *State* → *Action*) function
- This is a dynamic (time-sequenced control) system under uncertainty

Formal MDP Framework

The following notation is for discrete time steps. Continuous-time formulation is analogous (often involving <u>Stochastic Calculus</u>)

- Time steps denoted as $t = 1, 2, 3, \ldots$
- Markov States $S_t \in \mathcal{S}$ where $\mathcal{S}$ is the State Space
- Actions $A_t \in \mathcal{A}$ where $\mathcal{A}$ is the Action Space
- Rewards $R_t \in \mathbb{R}$ denoting numerical feedback
- Transitions $p(s', r | s, a) = Pr\{S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a\}$
- $\gamma \in [0, 1]$ is the Discount Factor for Reward when defining *Return*
- Return $G_t = R_t + \gamma \cdot R_{t+1} + \gamma^2 \cdot R_{t+1} + \ldots$
- Policy $\pi(a|s)$ is probability that Agent takes action a in states s
- The goal is find a policy that maximizes $\mathbb{E}[G_t | S_t = s]$ for all $s \in \mathcal{S}$

Positive/negative feedback

Posture, orientation

Baby steps

World

Many real-world problems fit this MDP framework

- Self-driving vehicle (speed/steering to optimize safety/time)
- Game of Chess (Boolean *Reward* at end of game)
- Complex Logistical Operations (eg: movements in a Warehouse)
- Make a humanoid robot walk/run on difficult terrains
- Manage an investment portfolio
- Control a power station
- Optimal decisions during a football game
- Strategy to win an election (high-complexity MDP)

Self-Driving Vehicle

Ride quality, arrival time

Location, velocity, sensor data

Steering, acceleration, braking

Traffic, signals, road conditions, routes

Why are these problems hard?

- *State* space can be large or complex (involving many variables)
- Sometimes, *Action* space is also large or complex
- No direct feedback on "correct" *Actions* (only feedback is *Reward*)
- Time-sequenced complexity (*Actions* influence future *States/Actions*)
- *Action*s can have delayed consequences (late *Reward*s)
- *Agent* often doesn't know the *Model* of the *Environment*
- "Model" refers to probabilities of state-transitions and rewards
- So, *Agent* has to learn the *Model* AND solve for the Optimal *Policy*
- *Agent Action*s need to tradeoff between "explore" and "exploit"

Value Function and Bellman Equations

- Value function (under policy π) $V_\pi(s) = \mathbb{E}[G_t | S_t = s]$ for all $s \in \mathcal{S}$

$$V_\pi(s) = \sum_a \pi(a|s) \sum_{s',r} p(s', r|s, a) \cdot (r + \gamma V_\pi(s')) \text{ for all } s \in \mathcal{S}$$

- Optimal Value Function $V_*(s) = \max_\pi V_\pi(s)$ for all $s \in \mathcal{S}$

$$V_*(s) = \max_a \sum_{s',r} p(s', r|s, a) \cdot (r + \gamma V_*(s')) \text{ for all } s \in \mathcal{S}$$

- *There exists an Optimal Policy π_* achieving $V_*(s)$ for all $s \in \mathcal{S}$*
- Determining $V_\pi(s)$ known as *Prediction*, and $V_*(s)$ known as *Control*
- The above recursive equations are called *Bellman equations*
- In continuous time, refered to as *Hamilton-Jacobi-Bellman (HJB)*
- The algorithms based on Bellman equations are broadly classified as:
  - Dynamic Programming
  - Reinforcement Learning

# Dynamic Programming versus Reinforcement Learning

- When Model is known $\Rightarrow$ *Dynamic Programming* (DP)
- DP Algorithms take advantage of knowledge of probabilities
- So, DP Algorithms do not require interaction with the environment
- Model-based/DP algorithms often refered to as *Planning Algorithms*
- When Model is unknown $\Rightarrow$ *Reinforcement Learning* (RL)
- RL Algorithms interact with the Environment and incrementally learn
- Environment interaction could be real interaction or a simulator
- RL approach: Try different actions & learn what works, what doesn't
- RL Algorithms' key challenge is to tradeoff "explore" versus "exploit"
- DP or RL, Good approximation of Value Function is vital to success
- Deep Neural Networks are typically used for function approximation

Why is RL interesting/useful to learn about?

- RL solves MDP problem when *Environment Model* is unknown
- Or when only an *Environment Simulator* is available
- The above two situations are typical in real-world problems
- **Promise of modern A.I. is based on success of RL algorithms**
- Potential for automated decision-making in many industries
- In 10-20 years: Bots that act or behave more optimal than humans
- RL already solves various low-complexity real-world problems
- RL might soon be the most-desired skill in the technical job-market
- Possibilities in Finance are endless (we cover 3 important problems)
- Learning RL is a lot of fun! (interesting in theory as well as coding)

Many Faces of Reinforcement Learning

Vague (but in-vogue) Classification of Machine Learning

Overview of the Course

- Theory of Markov Decision Processes (MDPs)
- Dynamic Programming (DP) Algorithms
- Reinforcement Learning (RL) Algorithms
- Plenty of Python implementations of models and algorithms
- Apply these algorithms to 3 Financial/Trading problems:
  - (Dynamic) Asset-Allocation to maximize Utility of Consumption
  - Optimal Exercise/Stopping of Path-dependent American Options
  - Optimal Trade Order Execution (managing Price Impact)
- By treating each of the problems as MDPs (i.e., Stochastic Control)
- We will go over classical/analytical solutions to these problems
- Then introduce real-world considerations, and tackle with RL (or DP)

Optimal Asset Allocation to Maximize Consumption Utility

- You can invest in (allocate wealth to) a collection of assets
- Investment horizon is a fixed length of time
- Each risky asset has an unknown distribution of returns
- Transaction Costs & Constraints on trading hours/quantities/shorting
- Allowed to consume a fraction of your wealth at specific times
- Dynamic Decision: Time-Sequenced Allocation & Consumption
- To maximize horizon-aggregated Utility of Consumption
- Utility function represents degree of risk-aversion
- So, we effectively maximize aggregate Risk-Adjusted Consumption

MDP for Optimal Asset Allocation problem

- *State* is [Current Time, Current Holdings, Current Prices]
- *Action* is [Allocation Quantities, Consumption Quantity]
- *Action*s limited by various real-world trading constraints
- *Reward* is Utility of Consumption less Transaction Costs
- *State*-transitions governed by risky asset movements

Optimal Exercise of Path-dependent American Options

- An American option can be exercised anytime before option maturity
- Key decision at any time is to exercise or continue
- The default algorithm is Backward Induction on a tree/grid
- But it doesn't work for path-dependent options
- Also, it's not feasible when state dimension is large
- Industry-Standard: Longstaff-Schwartz's simulation-based algorithm
- RL is an attractive alternative to Longstaff-Schwartz
- RL is straightforward once Optimal Exercise is modeled as an MDP

MDP for Optimal Options Exercise

- *State* is [Current Time, History of Underlying Security Prices]
- *Action* is Boolean: Exercise (i.e., Payoff and Stop) or Continue
- *Reward* always 0, except upon Exercise ($=$ Payoff)
- *State*-transitions governed by Underlying Prices' Stochastic Process
- Optimal Policy $\Rightarrow$ Optimal Stopping $\Rightarrow$ Option Price
- Can be generalized to other Optimal Stopping problems

Optimal Trade Order Execution (controlling Price Impact)

- You are tasked with selling a large qty of a (relatively less-liquid) stock
- You have a fixed horizon over which to complete the sale
- Goal is to maximize aggregate sales proceeds over horizon
- If you sell too fast, *Price Impact* will result in poor sales proceeds
- If you sell too slow, you risk running out of time
- We need to model temporary and permanent *Price Impact*s
- Objective should incorporate penalty for variance of sales proceeds
- Which is equivalent to maximizing aggregate Utility of sales proceeds

MDP for Optimal Trade Order Execution

- *State* is [Time Remaining, Stock Remaining to be Sold, Market Info]
- *Action* is Quantity of Stock to Sell at current time
- *Reward* is Utility of Sales Proceeds (i.e., Variance-adjusted-Proceeds)
- *Reward* & *State*-transitions governed by *Price Impact Model*
- Real-world *Model* can be quite complex (Order Book Dynamics)

Week by Week (Tentative) Schedule

- W1: Markov Decision Processes & Overview of Finance Problems
- W2: Bellman Equations & Dynamic Programming Algorithms
- W3: Optimal Asset Allocation problem
- W4: Optimal Exercise of American Options problem
- W5: Optimal Trade Order Execution problem, and Mid-Term Exam
- W6: Model-free Prediction (RL for Value Function Estimation)
- W7: Model-Free Control (RL for Optimal Value Function/Policy)
- W8: RL with Function Approximation (including Deep RL)
- W9: Batch Methods (DQN, LSTDQ/LSPI), and Gradient TD
- W10: Policy Gradient Algorithms
- W11: Final Exam

Sneak Peek into a few lectures in this course

- HJB Equation and Merton's Portfolio Problem
- Policy Gradient Theorem and Compatible Approximation Theorem
- Value Function Geometry and Gradient TD

Landmark Papers we will cover in detail

- Merton's solution for Optimal Portfolio Allocation/Consumption
- Longstaff-Schwartz Algorithm for Pricing American Options
- Bertsimas-Lo paper on Optimal Execution Cost
- Almgren-Chriss paper on Optimal Risk-Adjusted Execution Cost
- Original DQN paper and Nature DQN paper
- Lagoudakis-Parr paper on Least Squares Policy Iteration
- Sutton et al's paper on Policy Gradient

Similar Courses offered at Stanford

- AA 228/CS 238 (Mykel Kochenderfer - Autumn 2018)
- CS 234 (Emma Brunskill - Winter 2019)
- MS&E 251 (Edison Tse - Spring 2019)
- CS 332 (Emma Brunskill - Autumn 2018)
- MS&E 338 (Ben Van Roy - Spring 2019)
- MS&E 348 (Gerd Infanger - Winter 2020)
- MS&E 351 (Ben Van Roy - Winter 2019)